

training_logistic_regression

October 21, 2024

1 Implementing a Logistic Regression Classifier

1.1 1) Installing Libraries

```
[1]: # !pip install numpy pandas matplotlib --yes
```

```
[2]: #!pip install torch torchvision torchaudio
```

```
[3]: #!pip install watermark
```

```
[4]: #%load_ext watermark  
      #%watermark -v -p numpy,pandas,matplotlib,torch
```

1.2 2) Loading the Dataset

```
[5]: import pandas as pd  
  
df = pd.read_csv("perceptron_toydata-truncated.txt", sep="\t")  
df
```

```
[5]:
```

	x1	x2	label
0	0.77	-1.14	0
1	-0.33	1.44	0
2	0.91	-3.07	0
3	-0.37	-1.91	0
4	-0.63	-1.53	0
5	0.39	-1.99	0
6	-0.49	-2.74	0
7	-0.68	-1.52	0
8	-0.10	-3.43	0
9	-0.05	-1.95	0
10	3.88	0.65	1
11	0.73	2.97	1
12	0.83	3.94	1
13	1.59	1.25	1
14	1.14	3.91	1
15	1.73	2.80	1

```
16  1.31  1.85      1
17  1.56  3.85      1
18  1.23  2.54      1
19  1.33  2.03      1
```

```
[6]: X_train = df[["x1", "x2"]].values
     y_train = df["label"].values
```

```
[7]: X_train
```

```
[7]: array([[ 0.77, -1.14],
            [-0.33,  1.44],
            [ 0.91, -3.07],
            [-0.37, -1.91],
            [-0.63, -1.53],
            [ 0.39, -1.99],
            [-0.49, -2.74],
            [-0.68, -1.52],
            [-0.1 , -3.43],
            [-0.05, -1.95],
            [ 3.88,  0.65],
            [ 0.73,  2.97],
            [ 0.83,  3.94],
            [ 1.59,  1.25],
            [ 1.14,  3.91],
            [ 1.73,  2.8 ],
            [ 1.31,  1.85],
            [ 1.56,  3.85],
            [ 1.23,  2.54],
            [ 1.33,  2.03]])
```

```
[8]: X_train.shape
```

```
[8]: (20, 2)
```

```
[9]: y_train
```

```
[9]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1])
```

```
[10]: y_train.shape
```

```
[10]: (20,)
```

```
[11]: import numpy as np

      np.bincount(y_train)
```

```
[11]: array([10, 10])
```

1.3 3) Visualizing the dataset

```
[12]: %matplotlib inline
import matplotlib.pyplot as plt
```

```
[13]: plt.plot(
    X_train[y_train == 0, 0],
    X_train[y_train == 0, 1],
    marker="D",
    markersize=10,
    linestyle="",
    label="Class 0",
)


plt.plot(
    X_train[y_train == 1, 0],
    X_train[y_train == 1, 1],
    marker="^",
    markersize=13,
    linestyle="",
    label="Class 1",
)

plt.legend(loc=2)

plt.xlim([-5, 5])
plt.ylim([-5, 5])

plt.xlabel("Feature  $x_1$ ", fontsize=12)
plt.ylabel("Feature  $x_2$ ", fontsize=12)

plt.grid()
plt.show()
```



```
[14]: X_train = (X_train - X_train.mean(axis=0)) / X_train.std(axis=0)
```

```
[15]: plt.plot(
    X_train[y_train == 0, 0],
    X_train[y_train == 0, 1],
    marker="D",
    markersize=10,
    linestyle="",
    label="Class 0",
)


plt.plot(
    X_train[y_train == 1, 0],
    X_train[y_train == 1, 1],
    marker="^",
    markersize=13,
    linestyle="",
    label="Class 1",
)

plt.legend(loc=2)
```

```
plt.xlim([-5, 5])
plt.ylim([-5, 5])

plt.xlabel("Feature  $x_1$ ", fontsize=12)
plt.ylabel("Feature  $x_2$ ", fontsize=12)

plt.grid()
plt.show()
```



1.4 4) Implementing the model

```
[16]: import torch

class LogisticRegression(torch.nn.Module):

    def __init__(self, num_features):
        super().__init__()
        self.linear = torch.nn.Linear(num_features, 1)
```

```

def forward(self, x):
    logits = self.linear(x)
    probas = torch.sigmoid(logits)
    return probas

```

```

[17]: torch.manual_seed(1)

model = LogisticRegression(num_features=2)

```

```

[18]: x = torch.tensor([1.1, 2.1])

with torch.no_grad():
    proba = model(x)

print(proba)

```

```

tensor([0.4033])

```

1.5 5) Defining a DataLoader

```

[19]: from torch.utils.data import Dataset, DataLoader

class MyDataset(Dataset):
    def __init__(self, X, y):

        self.features = torch.tensor(X, dtype=torch.float32)
        self.labels = torch.tensor(y, dtype=torch.float32)

    def __getitem__(self, index):
        x = self.features[index]
        y = self.labels[index]
        return x, y

    def __len__(self):
        return self.labels.shape[0]

train_ds = MyDataset(X_train, y_train)

train_loader = DataLoader(
    dataset=train_ds,
    batch_size=10,
    shuffle=True,
)

```

```

[20]: X_train.shape

```

[20]: (20, 2)

1.6 6) The training loop

```
[21]: import torch.nn.functional as F

torch.manual_seed(1)
model = LogisticRegression(num_features=2)
optimizer = torch.optim.SGD(model.parameters(), lr=0.05)

num_epochs = 20

for epoch in range(num_epochs):

    model = model.train()
    for batch_idx, (features, class_labels) in enumerate(train_loader):

        probas = model(features)

        loss = F.binary_cross_entropy(probas, class_labels.view(probas.shape))

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    ### LOGGING
    print(f'Epoch: {epoch+1:03d}/{num_epochs:03d}'
          f' | Batch {batch_idx:03d}/{len(train_loader):03d}'
          f' | Loss: {loss:.2f}')
```

```
Epoch: 001/020 | Batch 000/002 | Loss: 0.67
Epoch: 001/020 | Batch 001/002 | Loss: 0.73
Epoch: 002/020 | Batch 000/002 | Loss: 0.67
Epoch: 002/020 | Batch 001/002 | Loss: 0.67
Epoch: 003/020 | Batch 000/002 | Loss: 0.60
Epoch: 003/020 | Batch 001/002 | Loss: 0.68
Epoch: 004/020 | Batch 000/002 | Loss: 0.69
Epoch: 004/020 | Batch 001/002 | Loss: 0.54
Epoch: 005/020 | Batch 000/002 | Loss: 0.61
Epoch: 005/020 | Batch 001/002 | Loss: 0.57
Epoch: 006/020 | Batch 000/002 | Loss: 0.59
Epoch: 006/020 | Batch 001/002 | Loss: 0.54
Epoch: 007/020 | Batch 000/002 | Loss: 0.51
Epoch: 007/020 | Batch 001/002 | Loss: 0.58
Epoch: 008/020 | Batch 000/002 | Loss: 0.51
```

```

Epoch: 008/020 | Batch 001/002 | Loss: 0.54
Epoch: 009/020 | Batch 000/002 | Loss: 0.51
Epoch: 009/020 | Batch 001/002 | Loss: 0.49
Epoch: 010/020 | Batch 000/002 | Loss: 0.53
Epoch: 010/020 | Batch 001/002 | Loss: 0.44
Epoch: 011/020 | Batch 000/002 | Loss: 0.42
Epoch: 011/020 | Batch 001/002 | Loss: 0.52
Epoch: 012/020 | Batch 000/002 | Loss: 0.46
Epoch: 012/020 | Batch 001/002 | Loss: 0.44
Epoch: 013/020 | Batch 000/002 | Loss: 0.46
Epoch: 013/020 | Batch 001/002 | Loss: 0.42
Epoch: 014/020 | Batch 000/002 | Loss: 0.41
Epoch: 014/020 | Batch 001/002 | Loss: 0.43
Epoch: 015/020 | Batch 000/002 | Loss: 0.43
Epoch: 015/020 | Batch 001/002 | Loss: 0.39
Epoch: 016/020 | Batch 000/002 | Loss: 0.41
Epoch: 016/020 | Batch 001/002 | Loss: 0.39
Epoch: 017/020 | Batch 000/002 | Loss: 0.38
Epoch: 017/020 | Batch 001/002 | Loss: 0.39
Epoch: 018/020 | Batch 000/002 | Loss: 0.44
Epoch: 018/020 | Batch 001/002 | Loss: 0.31
Epoch: 019/020 | Batch 000/002 | Loss: 0.34
Epoch: 019/020 | Batch 001/002 | Loss: 0.39
Epoch: 020/020 | Batch 000/002 | Loss: 0.33
Epoch: 020/020 | Batch 001/002 | Loss: 0.38

```

1.7 7) Evaluating the results

```
[22]: probas
```

```

[22]: tensor([[0.6687],
              [0.6810],
              [0.2162],
              [0.7398],
              [0.2126],
              [0.5744],
              [0.3209],
              [0.6503],
              [0.4276],
              [0.2598]], grad_fn=<SigmoidBackward0>)

```

```

[23]: pred = torch.where(probas > 0.5, 1, 0)
      pred

```

```

[23]: tensor([[1],
              [1],
              [0],

```



```
[1],  
[0],  
[1],  
[0],  
[1],  
[0],  
[0]])
```

```
[24]: class_labels.view(pred.shape).to(pred.dtype)
```

```
[24]: tensor([[1],  
            [1],  
            [0],  
            [1],  
            [0],  
            [1],  
            [0],  
            [1],  
            [0],  
            [0]])
```

```
[25]: def compute_accuracy(model, dataloader):  
  
    model = model.eval()  
  
    correct = 0.0  
    total_examples = 0  
  
    for idx, (features, class_labels) in enumerate(dataloader):  
  
        with torch.no_grad():  
            probas = model(features)  
  
            pred = torch.where(probas > 0.5, 1, 0)  
            lab = class_labels.view(pred.shape).to(pred.dtype)  
  
            compare = lab == pred  
            correct += torch.sum(compare)  
            total_examples += len(compare)  
  
    return correct / total_examples
```

```
[26]: train_acc = compute_accuracy(model, train_loader)
```

```
[27]: print(f"Accuracy: {train_acc*100}%")
```

Accuracy: 100.0%

1.8 8) Optional: visualizing the decision boundary

```
[28]: plt.plot(
        X_train[y_train == 0, 0],
        X_train[y_train == 0, 1],
        marker="D",
        markersize=10,
        linestyle="",
        label="Class 0",
    )


    plt.plot(
        X_train[y_train == 1, 0],
        X_train[y_train == 1, 1],
        marker="^",
        markersize=13,
        linestyle="",
        label="Class 1",
    )

    plt.legend(loc=2)

    plt.xlim([-5, 5])
    plt.ylim([-5, 5])

    plt.xlabel("Feature $x_1$", fontsize=12)
    plt.ylabel("Feature $x_2$", fontsize=12)

    plt.grid()
    plt.show()
```



```
[29]: def plot_boundary(model):

    w1 = model.linear.weight[0][0].detach()
    w2 = model.linear.weight[0][1].detach()
    b = model.linear.bias[0].detach()

    x1_min = -20
    x2_min = -(w1 * x1_min) - b / w2

    x1_max = 20
    x2_max = -(w1 * x1_max) - b / w2

    return x1_min, x1_max, x2_min, x2_max
```

```
[30]: x1_min, x1_max, x2_min, x2_max = plot_boundary(model)

plt.plot(
    X_train[y_train == 0, 0],
    X_train[y_train == 0, 1],
    marker="D",
```

```

        markersize=10,
        linestyle="",
        label="Class 0",
    )

plt.plot(
    X_train[y_train == 1, 0],
    X_train[y_train == 1, 1],
    marker="^",
    markersize=13,
    linestyle="",
    label="Class 1",
)

plt.plot([x1_min, x1_max], [x2_min, x2_max], color="k")

plt.legend(loc=2)

plt.xlim([-5, 5])
plt.ylim([-5, 5])

plt.xlabel("Feature $x_1$", fontsize=12)
plt.ylabel("Feature $x_2$", fontsize=12)

plt.grid()
plt.show()

```

